

AI DEMYSTIFIED

MUSCAT, OMAN

THE MASTER MINUTE-BY-MINUTE FACILITATOR SCRIPT

3 Days · 8 Hours per Day · 09:00 – 17:00

Instructors

Amardeep Bajpai — Business Architect
Atharv Kumar — Technical Wizard

Partners: Traininglobe & SIG Oman

Tone: High-energy, conversational, story-driven, authoritative.

How to Read This Script — The Engagement Layer

This script carries three running gags across all three days so the room builds real familiarity rather than hearing isolated jokes: “Ahmed from Accounts” (a good-natured stand-in for the un-upgraded employee, who gets a genuine redemption arc by the Day 3 close), “the Family WhatsApp Group” (the go-to example of glorious, chaotic, unstructured data), and “the Intern” (AI as brilliant-but-naive help that needs managing). Lines marked with 😊 are laugh beats — deliver the punchline, then pause and let the room actually laugh before moving on. Purple-shaded blocks are short Story Time anecdotes; slow down and tell these like a story, not a bullet list.

DAY 1: AI FUNDAMENTALS (Demystify & Empower)

Goal for the Day: Eradicate AI fear. Move them from "AI is magic/scary" to "AI is just math and a very fast intern." By 17:00, they will be writing advanced prompts.

09:00 – 09:05 Doors / Stage Silence

[Stage Direction] Screen is black or on a generic holding slide with soft instrumental music under 40%. Amardeep and Atharv stand offstage. House lights stay up until 08:58, then dim slowly. Amardeep walks to dead-center stage at 09:00 sharp. No mic check chatter — silence is the opener.

Facilitator Note: Facilitator cue: do not start until the room has genuinely quieted. A 10-15 second silence before the first word is deliberate — it resets the room from "conference chatter" to "performance."

09:05 – 09:15 The Hook — Hands Up

Amardeep: "Good morning, Muscat. Let me ask you a very honest question before we look at a single slide today. Raise your hand if you have used ChatGPT, Claude, or Gemini in the last 7 days."

[Stage Direction] Keep your own hand up. Wait a full 3 seconds for hands to rise. Scan the room left to right.

Amardeep: "Keep them up. Now, keep your hand raised if you feel, deep down, that you are using it to its absolute maximum potential."

[Stage Direction] Watch hands drop — most will. Smile. Let the silence do the work for 2-3 seconds.

Amardeep: "Exactly. Right now, most of the world is treating a multi-billion-dollar supercomputer like a glorified spelling checker. Or a fancy Google search. Over the next 24 hours of instruction across these three days, we are going to fix that."

😄 "And look, I get it. Somewhere in this room right now, there is an 'Ahmed from Accounts' who used ChatGPT once, in 2023, to write a birthday message for his mother-in-law, and has considered himself 'AI-trained' ever since." [Laugh Beat — pause, let it land, scan for the person actually nodding along]

Facilitator Note: Running gag planted: "Ahmed from Accounts" is introduced here as a shared, affectionate joke character. He will return on Day 2 and take a full redemption arc on Day 3 — do not skip this line, it is load-bearing for later punchlines.

Amardeep: "My name is Amardeep Bajpai, CEO of Webisdom. I look at AI and I see one thing: ROI. Scale. Efficiency. And joining me is Atharv Kumar. I build the business strategies; Atharv actually builds the brains. He looks at AI and sees neural weights, parameters, and Python code. Together, we are going to translate the code into commercial value for you."

Amardeep: "I also want to extend a massive thank you to our partners, Traininglobe and SIG Oman, for making this happen."

Facilitator Note: Timing buffer: if applause runs long, let it. This block can flex 09:05-09:17 without breaking the day.

09:15 – 09:30 The Mindset Shift

[Stage Direction] Click to Slide: "Artificial Intelligence Demystified — Title / The Shift"

Amardeep: "The media wants you terrified of AI. They tell you it's coming for your job. Write this phrase down on the notepads in front of you. Please, write it down: AI will not replace you. A person using AI will."

[Stage Direction] Pause 5 seconds — physically watch the room write. Do not talk over them.

Amardeep: "Look, we worry about AI taking over the world. But let's be real — AI doesn't know how to navigate the Muscat Expressway at 5 PM. AI doesn't understand your company's internal politics. AI cannot read a room of investors. It lacks human context. You have the context. AI is just your new intern."

It is a brilliant, hyper-fast, but occasionally very naive intern. Your job is no longer to do the busy work; your job is to manage the intern."

😄 "Think about your actual worst intern ever. Brilliant on paper, calls the client by the wrong name, and once replied-all to the entire company. That is AI without supervision. Charming, fast, occasionally catastrophic. You are not being replaced — you are being promoted to 'intern manager.'" [Laugh Beat]

Facilitator Note: Optional meme insert here (per raw notes): a slide showing a stressed office worker vs. a relaxed one "managing the intern." Use only if the crowd is still settling — otherwise skip straight to 09:30.

09:30 – 09:45 The Intelligence Ladder, Part 1: Data → Information

[Stage Direction] Click to Slide: "The Intelligence Ladder" — Video/Animation cue #1 plays here (60-90 sec) before Atharv speaks.

Amardeep: "Atharv, let's break down what AI actually is, because people throw the word around like magic."

Atharv: "Right. It's not magic. It's a staircase. Let's look at the screen. At the bottom, we have Data. Imagine I hand you a piece of paper with the number '45' on it. It's meaningless."

Amardeep: "Completely useless."

Atharv: "Exactly. So we move up a step to Information. I add a label. It now says, 'Temperature in Muscat: 45 Degrees.' Now it's a fact. But a fact doesn't make you money."

😄 Amardeep: "For a real-world example of pure, uncut Data with zero structure — open your family WhatsApp group right now. Voice notes, forwarded videos, seventeen good-morning GIFs before 7 AM. That is Data at its most chaotic. Somebody, somewhere, needs to build an AI just to summarize what actually happened in that group today." [Laugh Beat — this usually gets the biggest laugh of the morning; let it breathe]

Facilitator Note: Running gag planted: "The Family WhatsApp Group" as the go-to example of unstructured data. Callback this on Day 2 when discussing training data and hallucinations.

09:45 – 10:00 The Intelligence Ladder, Part 2: Knowledge → Wisdom → AI

Amardeep: "So we move to Knowledge."

Atharv: "Knowledge is historical context. We look at our sales data and realize: 'Whenever it hits 45 degrees, Karak tea sales drop by 40%, but iced coffee sales spike by 200%.'"

Amardeep: "Then comes the human element. Wisdom. The manager says, 'Okay, it's going to be 45 degrees tomorrow. Stop the Karak ads, double the budget on the iced coffee ads.'"

Amardeep: "And here is where the world changes. For the last 50 years, computers stopped at Information. Humans had to do Knowledge and Wisdom. Today, we step up to AI. AI is the autopilot. AI is the system that connects to the global weather API, sees the forecast is 45 degrees, automatically pauses the Karak Facebook ads, automatically shifts the budget to Iced Coffee, and emails the warehouse to stock up on ice — all at 3:00 AM while you are asleep. That is the shift from manual to autonomous."

10:00 – 10:15 ML vs DL vs AI — Video Cue + Quick Poll

[Stage Direction] Video/Animation cue #2: "ML vs DL vs AI" plays (90 sec). Follow immediately with a live poll (show of hands or Slido/Mentimeter): "Which of these three terms do you use most often — and did you actually mean it?"

Facilitator Note: Keep the poll to 3 minutes max. This is a warm-up for the deeper ML/DL/AI session at 11:30 on Day 1 — don't over-explain here, just plant the question.

10:15 – 10:30 ANI, AGI, ASI — Video Cue + Framing

[Stage Direction] Video/Animation cue #3: "ANI / AGI / ASI" plays (90 sec).

Amardeep: "Hollywood lied to us. The Matrix, Skynet, Ex Machina — that's ASI, Artificial Super Intelligence, vastly smarter than all of humanity combined. We are decades away from that. Right now we live in the era of ANI — Artificial Narrow AI. Master of one task, useless outside it. Our goal this week is not to fear the future, but to master the narrow tools already in front of us."

10:30 – 10:45 Q&A and Interactive Discussion

Amardeep: "Before we take our morning coffee, I want to hear from you. Who in this room spends more than 3 hours a day just gathering Data and Information? Making Excel sheets match? Formatting reports?"

[Stage Direction] Engage 2-3 audience members directly. Ask their names, their roles, and validate their pain points on mic.

Amardeep: "By Day 3, we are going to automate that exact bottleneck for you."

10:45 – 11:00 Topic 4 Teaser — Master Research Through AI / Prompt Engineering Seed

[Stage Direction] Video/Animation cue #4 plays under soft background music — a fast-cut teaser of what prompt engineering can do (research, writing, images). This is a deliberate appetite-builder for the 14:00 Prompt Engineering session.

Amardeep: "Hold that thought — at 2 PM today, we're going to hand you the exact framework the top 1% use to talk to these models. Go get your coffee."

Facilitator Note: 11:00 – 11:30 | COFFEE & NETWORKING BREAK (30 min). Music playlist: upbeat instrumental. Facilitators circulate informally — this is unscripted relationship-building time, not downtime.

11:30 – 11:45 ML vs DL vs AI — Concentric Circles

[Stage Direction] Click to Slide: "AI > ML > DL"

Amardeep: "Welcome back. If you want to command respect in a boardroom, you need to use the right words. AI, Machine Learning, and Deep Learning are not synonyms."

Atharv: "Let's use an analogy. AI is the Universe. Machine Learning is a Galaxy inside it. Deep Learning is our Solar System."

11:45 – 12:00 Machine Learning — The Netflix Example

Atharv: "Let's talk about Machine Learning. ML is learning without being explicitly programmed. When you open Netflix tonight, how does it know to recommend a specific movie? No programmer wrote a line of code saying, 'If Ahmed likes Batman, show him Iron Man.' That would take forever. Instead, the ML algorithm noticed that 10,000 people in the Middle East who watched Batman also watched Iron Man. It learned from the pattern."

😊 **Amardeep:** "And here's where it gets personal — you watch one documentary about a dangerous animal at 1 AM, and suddenly your entire homepage thinks you're planning something. Your wife looks over your shoulder at the recommendations and now you have some explaining to do. That's the algorithm. It's not judging you. It's just very, very confident." [Laugh Beat]

Facilitator Note: Optional POLL/QUIZ slot here: "Name one Netflix/YouTube/TikTok recommendation that felt eerily accurate." Take 2 verbal answers, keep it under 3 minutes.

12:00 – 12:15 Deep Learning — The Tesla Example

Amardeep: "And Deep Learning? Why is it 'Deep'?"

Atharv: "Because it uses layers of neural networks to process chaotic, unstructured data — vision and voice. When a Tesla on Autopilot sees a shadow, a plastic bag blowing in the wind, and a child running into the street, it has to know instantly which one to brake for. That level of complex, layered processing is Deep Learning."

12:15 – 12:30 ANI vs AGI vs ASI — The Reality Check

[Stage Direction] Click to Slide: "Timeline of AI"

Amardeep: "We need to ground our expectations. Right now, we live in the era of ANI — Artificial Narrow AI. ChatGPT might write a brilliant legal contract in 3 seconds, but if you ask it to make a cup of tea or fix a leaky pipe, it is useless. It is a 'Master of One.'"

12:30 – 12:45 Interactive Puzzle / Quiz — Classify the AI

[Stage Direction] Run a rapid-fire quiz: show 5 real-world examples (a chess engine, a self-driving car, a spam filter, a hypothetical superintelligence, a chatbot) and have the audience shout ANI, AGI, or ASI for each. Keep pace fast — under 10 minutes.

12:45 – 13:00 Session Wrap + Small Awards Teaser

Amardeep: "Great energy. Hold onto that — we hand out our first spot prizes at the end of today. Let's break for lunch and come back sharp, because right after lunch is the single most useful 90 minutes of the entire three days: Prompt Engineering."

Facilitator Note: 13:00 – 14:00 | LUNCH BREAK (60 min).

14:00 – 14:10 The Post-Lunch Wake-Up

[Stage Direction] Audience will be sleepy. Keep energy incredibly high. Walk into the aisles, use a wireless mic if available.

Amardeep: "Alright, I know you just had lunch. But this next 90 minutes is the most important session of Day 1. If you learn nothing else today, learn this."

14:10 – 14:25 The CREATE Framework — Introduced

Atharv: "We are going to learn how to talk to the machine. Let's look at the CREATE Framework. Most people treat AI like Google. They type: 'Tell me about retail in Oman.' And the AI gives them a Wikipedia-style, boring answer. CREATE stands for Context, Request, Explanation, Action, Tone, and Extras. You must give the AI a persona. 'Act as a Senior Market Research Analyst.' You must give it a format. 'Give me a Markdown table.' You must give it a tone. 'Keep it persuasive and data-driven.' If you don't dictate these terms, the AI defaults to average."

Facilitator Note: Video/Animation cue #5 plays here under background music — a visual breakdown of C-R-E-A-T-E, one letter lighting up at a time (approx. 60 sec).

14:25 – 14:35 Live Demo — The Bad Prompt (a.k.a. "Roast the Prompt")

[Stage Direction] Atharv sits at the laptop, mirroring screen to projector. Opens ChatGPT.

Atharv: Types: "Write an email to my boss asking for a bigger budget for software." Hits enter. Shows the generic response.

Amardeep: "Look at that. 'Dear Boss, I hope this email finds you well...' Nobody talks like that. It sounds robotic."

😄 "This is what we call an 'Ahmed from Accounts' prompt. Vague, polite, and about as persuasive as a parking ticket. Everybody, on the count of three, read the fake email out loud in your most robotic voice... one, two, three—" [Laugh Beat — let the room chorus it badly and laugh at themselves]

Facilitator Note: Facilitator tip: this "read it in your robot voice" moment is a cheap, reliable laugh and a physical energy reset after 25 minutes of sitting. Don't skip it even if running slightly behind.

14:35 – 14:50 Live Demo — The CREATE Prompt

Amardeep: "Now, Atharv, use the CREATE framework."

[Stage Direction] Atharv types: "Act as a persuasive B2B sales director (Context). Write a 3-paragraph email to my CEO (Request). I need \$50,000 for a new CRM system because our current Excel sheets are causing us to lose 15% of inbound leads (Explanation). Use bullet points for the ROI (Action). Keep the tone urgent but highly professional, sounding like a seasoned executive (Tone)."

Amardeep: "Boom. Look at the difference. It's quantifying the loss. It's structuring the argument. It sounds like a human executive wrote it. This is prompt engineering."

14:50 – 15:05 Breakdown & Class Discussion

[Stage Direction] Take 2-3 audience examples verbally: "Give me a task you do over and over — let's CREATE-ify it live." Type one on the spot if time allows.

15:05 – 15:30 Bridge to Hands-On + Rules of the Hackathon

Amardeep: "Here's what happens after the break. Laptops open, no more listening to us — you build a real prompt for a real problem sitting on your desk right now. Best prompts win a prize."

Facilitator Note: 15:30 – 15:45 | AFTERNOON BREAK (15 min).

15:45 – 16:00 The Prompt Hackathon — Setup

Amardeep: "Laptops open. You have 45 minutes. Take a real problem sitting on your desk right now — an email you've been putting off, a report you need to summarize, a strategy you need to draft. Build a CREATE prompt for it."

[Stage Direction] Post the CREATE checklist on screen for reference throughout the exercise.

16:00 – 16:30 The Prompt Hackathon — Working Time

[Stage Direction] Atharv and Amardeep walk the floor constantly, looking over screens. Point out good work loudly to the room.

Amardeep: "Look at what Tariq is doing over here — he fed the AI his company's actual brand guidelines before asking it to write! Brilliant."

😄 Amardeep (mock-horrified, walking past a lazy prompt): "Oh no. Oh no no no. This is a crime scene. Someone typed 'help me' with no context and hit enter. That's not a prompt, that's a cry for help. We'll come back and rescue this one before the 45 minutes are up." [Laugh Beat]

Facilitator Note: Facilitator tip: aim to publicly praise at least 4-5 different participants by name during this block — spread recognition across the room, not just the front rows. The mock-horror bit above works best used sparingly, once or twice, on a genuinely good-natured target who can take the joke.

16:30 – 16:45 The Prompt Hackathon — Final Stretch & Volunteer Shares

[Stage Direction] Invite 2-3 volunteers to read their prompt and result aloud on mic.

16:45 – 17:00 Day 1 Wrap-Up, Spot Awards & Seeding the Pitch

Amardeep: "Time is up! Tariq and Fatima, please come up here. You two had the best prompts of the day. Here is a small token from Webisdom."

[Stage Direction] Hand out chocolates/swag on stage.

Amardeep: "Tonight, do not just close your laptops. Go home and show this to your families. Use it to plan a vacation itinerary. Tomorrow at 09:00 sharp, we open the hood of the car. We are going to see exactly how this machine thinks. Have a great evening!"

😊 "And somewhere out there tonight, Ahmed from Accounts is going to open ChatGPT for the second time in three years. Let him have his moment." [Laugh Beat — warm, not mocking; sets up his Day 3 comeback]

Amardeep: "One more thing before you go — everything we cover these three days is Foundation level. There's a Practitioner level beyond this, where you stop learning about AI and start building with AI. You'll see exactly what that looks like by Day 3."

DAY 2: AI FRAMEWORKS (Under the Hood)

Goal for the Day: Shift them from "users" to "mechanics." Understand Neural Networks, Transformers, Reinforcement Learning, hallucinations, and Visual AI — plus the risks that come with all of it.

09:00 – 09:15 Welcome Back & Recap

Amardeep: "Good morning! Yesterday we learned how to drive the car. Today, we are opening the engine block. Because if you don't know how the engine works, you won't know how to fix it when it breaks — and AI will break. It will hallucinate. Atharv, take us into the matrix."

09:15 – 09:35 How AI "Sees" — Neural Networks

[Stage Direction] Click to Slide: Neural Networks / Image of a Cat. Video/Animation cue plays showing layers activating on the cat image (approx. 90 sec).

Atharv: "Look at this cat on the screen. How do you know it's a cat? You see ears, fur, eyes. It takes your brain a millisecond. But a computer? It has no eyes. It only sees a grid of pixels. It sees numbers: 217, 154, 82."

Atharv: "A neural network is just a giant mathematical filter. It passes these numbers through 'layers.' The first layer looks for straight lines. The next layer looks for curves. The next layer looks for shapes. Finally, it outputs a probability: 'I am 98% sure this mathematical arrangement of pixels is a cat.'"

[Stage Direction] Click to a second image: side-by-side photos of a chihuahua and a blueberry muffin, genuinely hard to tell apart at a glance.

😄 **Amardeep:** "And this is exactly why early image models famously got confused between chihuahuas and blueberry muffins. Look at your screen — round, dark spots, roughly the same size. To a human, obviously different. To a machine looking purely at pixel patterns? A coin flip. This is why we test these things before we trust them with anything important." [Laugh Beat]

09:35 – 10:00 Why This Matters — Pattern Matching, Not Understanding

Atharv: "Why does this matter to you? Because AI does not understand what a cat is. It does not understand love, or business, or anger. It is purely a pattern-matching engine. If you remember that, you will never be fooled by it."

STORY TIME — THE LAWYER AND THE IMAGINARY COURT CASES

Amardeep: "Let me tell you a story that actually happened, more than once, in more than one country. A lawyer — a real, licensed, experienced lawyer — used an AI chatbot to help write a legal filing. He asked it to find supporting case law. The AI, very confidently, invented several court cases. Fake names, fake rulings, fake page numbers. It sounded completely real. The lawyer, trusting the confident tone, filed it with the court. The judge's clerk went looking for those cases. They did not exist. Not one of them. The lawyer was publicly embarrassed and sanctioned. Why do I tell you this on Day 2, and not Day 1? Because now you understand why it happened — the model is a pattern-matching engine trained to sound confident, not a database that checks its own facts. It doesn't know the difference between 'this is true' and 'this sounds like the kind of thing that is usually true.' Confidence is not the same as correctness. Never trust the tone. Always verify the source."

Facilitator Note: This is intentionally told as a general, widely-known cautionary pattern rather than naming a specific case or jurisdiction — the lesson lands the same either way, and it keeps the story evergreen regardless of which real-world incident the audience may have separately heard about.

Facilitator Note: POLL/QUIZ slot: show 3 optical-illusion-style images and ask the room what a neural net might misclassify and why. Ties the abstract concept to something visual and fun.

10:00 – 10:20 Transformers, Part 1 — "Attention Is All You Need"

Amardeep: "Before 2017, talking to an AI was awful. It had no memory. By the end of a sentence, it forgot how the sentence started. Then, researchers at Google published a paper. It had a very catchy title: 'Attention Is All You Need.'"

10:20 – 10:40 Transformers, Part 2 — The Lawyer Analogy

Atharv: "That paper introduced the 'Transformer' architecture. Think about how a lawyer reads a 50-page contract. The lawyer does not read the words 'the,' 'and,' 'it' with heavy focus. The lawyer pays attention to names, dates, and money. Transformers allowed AI to do exactly this. Instead of reading word-by-word, it looks at the whole paragraph at once and weighs the relationship between the words. It pays 'attention' to the context. That is why ChatGPT feels like it actually understands you."

[Stage Direction] Video/Animation cue plays showing an "attention map" lighting up connections between words in a sentence.

10:40 – 11:00 Interactive Puzzle — Spot the Attention

[Stage Direction] Show an ambiguous sentence on screen (e.g., "The trophy didn't fit in the suitcase because it was too big.") and ask the audience: what does "it" refer to? Use this to demonstrate why context-weighting matters. Keep to 15-20 minutes including discussion.

Facilitator Note: 11:00 – 11:30 | COFFEE BREAK (30 min).

11:30 – 11:50 RLHF — Why AI Is Polite

Atharv: "If we trained AI on the entire internet — including Twitter, Reddit, and all the dark corners of the web — why isn't ChatGPT completely insane, rude, and offensive? The answer is RLHF: Reinforcement Learning from Human Feedback."

Amardeep: "Atharv, how do you train a dog?"

Atharv: "You say 'Sit.' The dog barks. You do nothing. You say 'Sit.' The dog sits. You give it a treat. We did the exact same thing to AI. We had thousands of humans sit in rooms reading AI answers. If the AI gave a helpful, safe answer, the human clicked 'Thumbs Up' (+1 point). If it was rude or dangerous, 'Thumbs Down' (-1 point). We literally conditioned its behavior mathematically."

😊 **Amardeep:** "So the next time you type 'thanks, that's perfect!' to your AI, understand — you are not being polite for no reason. You are literally dropping a treat. You are training it in real time, every single conversation. Somewhere, a very well-behaved digital dog is doing tricks because of you." [Laugh Beat]

11:50 – 12:00 Reinforcement Learning — Quick Video Reinforcement

[Stage Direction] Short video/animation showing a simple RL loop (agent, environment, reward) — ties back to the dog-training analogy just given, approx. 60-90 sec.

12:00 – 12:10 Warm-Up Game — "AI or Human?"

[Stage Direction] Show 3 short text passages on screen, one at a time (a paragraph of real human writing, a paragraph of AI writing, and one blended/edited passage). Audience votes AI or Human by show of hands for each. Reveal the answer after each vote with a quick reason why it fooled or didn't fool the room.

😊 **Amardeep:** "Careful — I've watched entire rooms confidently vote 'definitely human' on something a machine wrote in four seconds. Don't feel bad if you get all three wrong. That's actually the whole point of today." [Laugh Beat]

Facilitator Note: This game doubles as a natural bridge into the LLM Showdown — it warms the room up to critically evaluate AI output before watching three models go head-to-head.

12:10 – 12:20 The Live LLM Showdown — Setup

[Stage Direction] Atharv hooks up laptop. Opens ChatGPT, Claude, and Gemini in three side-by-side windows.

Amardeep: "Everyone says 'ChatGPT' like it's the only AI. It's not. There is a war going on right now. OpenAI has ChatGPT, Google has Gemini, Anthropic has Claude. We are going to do a live battle. Atharv, give them a complex task."

😊 "Think of this next 15 minutes as our own little singing competition — except instead of judges with a buzzer, we have you. Three contestants. One task. Let's see who takes home the imaginary trophy." [Laugh Beat]

[Stage Direction] Atharv types into all three at once: "Write a Python script to scrape a real estate website for housing prices in Muscat, format it into a clean table, and write a 1-paragraph summary explaining this code to a non-technical CEO."

12:20 – 12:45 The Live LLM Showdown — Play-by-Play & Vote

Amardeep: (Doing play-by-play commentary like a sports announcer) "Look at ChatGPT go, the code is flying out. But wait, look at Claude — the formatting on Claude's explanation for the CEO is beautiful, very easy to read. And Gemini is actually pulling live data from Google Search."

Amardeep: "Let's vote! By round of applause, who thinks ChatGPT won? ... Claude? ... Gemini? The lesson here is: Don't be loyal to one tool. Use Claude for writing documents. Use ChatGPT for coding and logic. Use Gemini when you need live internet research."

😊 Amardeep: "And no, before anyone asks — none of the three models get a participation trophy. This isn't your child's football tournament. There is a real winner here, and it changes depending on the task, which is exactly the lesson." [Laugh Beat]

12:45 – 13:00 Bridge to Afternoon

Amardeep: "After lunch, we shift gears completely — from text to visuals, and then into a topic every single one of you needs to hear: what can actually go wrong."

Facilitator Note: 13:00 – 14:00 | LUNCH BREAK (60 min).

14:00 – 14:20 Directing the AI — Image Generation Formula

Amardeep: "Text saves time, but visuals sell. Whether it's marketing, presentations, or product design. But prompting an image is different than text. You are no longer a writer; you are a movie director."

Atharv: "The formula is: Subject + Action + Environment + Lighting + Camera Angle + Style."

14:20 – 14:40 Live Image Demo — Audience Suggestion

Atharv: "Let's do it live. Audience, shout out a crazy concept!"

[Stage Direction] Audience yells ideas. Example used: an Omani man on Mars. Atharv types into Midjourney/DALL-E: "An Omani astronaut wearing a traditional dishdasha under a spacesuit, drinking Karak tea on the surface of Mars, futuristic cyberpunk city in the background, cinematic lighting, 8k resolution, photorealistic." Wait 10 seconds. Show the stunning image. The room will react.

Amardeep: "Imagine paying a graphic designer to make that. It would take a week and cost \$500. We just did it in 10 seconds for fractions of a penny."

😊 Atharv: "Now, zoom in on the hand." [Zoom in on the astronaut's glove — if it has six fingers or a strange extra joint, point at it] "There it is. Every single time. AI is a brilliant artist and an absolutely terrible hand model. If you ever want to spot a fake image fast, count the fingers first." [Laugh Beat]

Facilitator Note: Repeat once more with a second audience suggestion if energy is high and time allows (up to 14:40). Also demo a short AI video-generation clip here per the raw agenda's "Image Generation + Video Generation" note.

14:40 – 15:05 The Dark Side — Deepfakes

[Stage Direction] Amardeep shifts to a very serious, hushed tone.

Amardeep: "Now, let's sober up. We just had a lot of fun. But look at that image on the screen. It looks 100% real. Now imagine a hacker doing that with your CEO's face. Or cloning your CFO's voice. Imagine getting a WhatsApp voice note from your boss, sounding exactly like him, saying, 'I'm stuck at customs, please wire \$10,000 to this account immediately.' This is happening right now."

15:05 – 15:30 Cybersecurity & Ethics/Governance — The Golden Rule

Atharv: "We call this the 'Zero Trust' era. You can no longer trust your eyes or your ears on a screen. You must verify through a secondary channel."

Amardeep: "And here is the Golden Rule of corporate AI: Never, ever paste confidential client data, financial sheets, or proprietary code into the free version of ChatGPT. You are feeding their global training model. If you do that, your competitor can ask ChatGPT a question tomorrow and it might spit out your financial data. If you need privacy, you use Enterprise versions, or you build local models."

Facilitator Note: This block also covers Session 9 from the raw agenda ("Ethics and Governance in AI") — frame governance as the shield that protects everything built in this workshop, not a separate compliance lecture.

Facilitator Note: 15:30 – 15:45 | AFTERNOON BREAK (15 min).

15:45 – 16:45 The "Ask Me Anything" Deep Dive

Amardeep: "We've covered heavy technical ground today. We are opening the floor. No question is too basic, no question is too complex. If it's business, I'll take it. If it's technical, Atharv will take it."

[Stage Direction] Spend 60 minutes doing rapid-fire Q&A. Use this time to continuously pivot their specific questions back to the idea that they need "structured training" and "automation," subtly setting up the Day 3 pitch.

Facilitator Note: Facilitator tip: keep a running mental tally of recurring pain points raised — reference 2-3 of them by name during the Day 3 pitch for a "we already know your problem" effect.

😊 **Amardeep** (if the "will AI take my job?" question comes up, which it always does): "I'll answer that honestly. AI will not take your job. But I promise you — Ahmed from Accounts, three floors down, who actually learned to use it properly? He might take yours." [Laugh Beat — this is the Day 2 payoff of the Ahmed gag; deliver it warmly, not as a threat, and let the room laugh before moving on]

16:45 – 17:00 Day 2 Wrap-Up

Amardeep: "You now know more about the mechanics of AI than 95% of your colleagues. But tomorrow is where we make the money. Tomorrow, we stop chatting, and we start automating. See you at 09:00."

Amardeep: "Foundation gets you fluent. Practitioner gets you employable-with-AI. More on that tomorrow."

DAY 3: AI FORMULATION & AUTOMATION (The Builder's Path)

Goal for the Day: Show massive workflow automation (APIs, RAG, Custom GPTs). Build up to the crescendo where they realize they need the Practitioner Program to execute this at scale.

Facilitator Note: FACILITATOR NOTE ON STRUCTURE: Two running orders were drafted for Day 3. Option 2 (tighter, hands-on-first) is the recommended default and is scripted in full below. Option 1 (topic-by-topic with a video/quiz per topic) is included afterward as a swappable alternate sequence — use it if the room responds better to shorter topic blocks than to one long automation build.

09:00 – 09:20 Day 3 Opening — The AI Transformation Framework

Amardeep: "Welcome to Day 3. The final day. For two days, you've been 'chatting' with AI. You type a prompt, it gives you an answer. That is great for personal productivity. But it doesn't scale a business. You can't hire a team just to sit and type prompts all day. Today, we move from adoption to automation."

Amardeep: "The question I get asked most is, 'Amardeep, where should we use AI in our company?' It's the wrong question. The right question is: 'What is the most painful, repetitive, soul-crushing bottleneck in your department?' Find that, and apply AI there."

09:20 – 09:35 AI Strategy — Framing the Day

Facilitator Note: This is Topic 10 from the raw agenda ("AI Strategy"), folded directly into the opening framework rather than run as a separate lecture — it is the "why" that the rest of the day builds toward.

Amardeep: "Every company in this room already has an AI strategy — it's just usually an accidental one, made up of whatever tools individual employees quietly adopted. By the end of today you'll have the outline of a deliberate one."

🤔 "Case in point: Ahmed from Accounts has quietly been using AI to draft his weekly reports for six months and never told his manager. That is not a strategy. That is a secret. Today we turn secrets into strategy." [Laugh Beat]

09:35 – 10:00 Functional AI & APIs — The Waiter Analogy

Facilitator Note: This block covers Topic 11 ("Functional AI") through the API explanation.

Atharv: "How do we get AI to work without us typing? We use APIs. What is an API? Imagine you go to a restaurant. You are the customer (the App). The kitchen is the server. You don't walk into the kitchen to cook your food. You give your order to the waiter. The waiter takes your request to the kitchen, gets the food, and brings it back. The waiter is the API. It connects two systems. Platforms like Zapier and Make.com are the best digital waiters in the world."

10:00 – 10:30 Workflow Automation Part 1 — Live Blueprint (Hands-On)

[Stage Direction] Click to slide showing a flowchart. Background music plays softly under this section per the raw agenda.

Amardeep: "Let's look at a live automation blueprint. Here is a process we build all the time. Trigger: a new lead fills out a form on your website. Action 1 (API): Zapier catches it and sends the data to OpenAI. AI Brain: OpenAI reads the lead's company size and asks, 'Is this a high-value prospect?' Action 2 (API): if yes, Zapier pushes the lead into Salesforce, drafts a customized welcome email, and sends a Slack notification to your top sales rep saying, 'Call this guy in the next 5 minutes.' Nobody touched a keyboard. This workflow runs 500 times a day while you are playing golf. That is automation."

[Stage Direction] Quick win: have participants sketch one automation trigger→action pair for their own job on paper or in a shared doc — this is the "hands-on, quick win" element from Option 2.

10:30 – 11:00 Ethics, Governance & Cybersecurity — "What Protects What You Just Built"

Amardeep: "Now that you've seen how much power an automated workflow has, we need to talk about what protects it. This is a combined session — ethics, governance, and cybersecurity — kept tight, because all three answer the same question: what happens when this system is wrong, biased, or attacked?"

Atharv: "Governance means: who is accountable when the AI makes a bad call in that Salesforce workflow? Ethics means: is the AI's scoring of 'high-value prospect' quietly biased against certain company types or regions? Cybersecurity means: if someone compromises your Zapier account, they now control the trigger, the AI brain, and the action, all at once. You don't bolt security on at the end — you design the trigger→AI→action chain with a human checkpoint before anything expensive or public happens."

Facilitator Note: This block satisfies Topics 12 and 13 from the raw agenda (Ethics/Governance, Cybersecurity) as one tight, framed session rather than two separate lectures.

Facilitator Note: 11:00 – 11:30 | COFFEE BREAK (30 min).

11:30 – 12:00 Introduction to Other AI Tools

Facilitator Note: Topic 14. Live-tour 4-5 tools beyond the "big three" — e.g. Perplexity for research, NotebookLM for document synthesis, Zapier/Make for automation (already seen), and one industry-specific tool relevant to the audience (retail/real estate/facilities, matching the Village Market case study to come).

Amardeep: "ChatGPT, Claude, and Gemini are the household names, but the tool that saves you the most hours is often one you've never heard of. Let's do a fast tour."

12:00 – 12:20 Training Your Own AI — Custom GPTs & RAG

Atharv: "ChatGPT is brilliant, but it has a major flaw: it doesn't know your business. It doesn't know your HR leave policy. It doesn't know the exact specs of the machines you sell. How do we fix that without spending millions of dollars training a new model? We use a technique called RAG: Retrieval-Augmented Generation. In simple terms, think of it as an 'Open Book Test' for the AI."

Amardeep: "You can build a 'Custom GPT.' You upload all 50 of your company's product manuals (PDFs) into the system. You write a system prompt: 'You are our Senior Tech Support Agent. Whenever a customer asks a question, look ONLY in these documents. If the answer isn't there, say you don't know. Do not hallucinate.' You just hired a perfect, 24/7 support agent in 20 minutes."

12:20 – 13:00 The Village Market Case Study

[Stage Direction] Click through the Village Market Case Study slides.

Amardeep: "Let's apply everything to a real case study. The Village Market. A massive commercial space. What are their pain points?"

STORY TIME — POOR ALI AND THE 1,000 SCREENS

Atharv: "Picture Ali, the night security guard. One human being, one small desk, and 1,000 camera feeds flickering in front of him like the world's worst multiplex. By 3 AM his eyes are doing that thing where they're open but nothing is actually registering. Somewhere on screen 640, a spill is forming into a hazard. Somewhere on screen 812, a bag has been sitting alone for eleven minutes. Ali cannot possibly catch both. Ali needs six more eyes, a double espresso, and a raise he's never going to get." [Laugh Beat] "That's the pain point. Now watch what Deep Learning does to Ali's night."

Atharv: "Pain point 1: Security. They have 1,000+ cameras. Humans physically cannot watch 1,000 screens. We deploy Deep Learning computer vision. The AI only alerts the guard if it sees a specific anomaly — a spill on the floor, or an unattended bag left for more than 5 minutes."

😄 Amardeep: "Suddenly Ali isn't watching 1,000 screens anymore. He's watching one notification that says 'Screen 640, possible spill, aisle 12.' Ali becomes a superhero. Well-rested, laser-focused, and still no raise — but let's not worry about that part today." [Laugh Beat]

Amardeep: "Pain point 2: Lease Management. They have hundreds of tenants, and leases are tracked on massive Excel sheets. Humans miss renewal dates. We build a workflow where the AI reads the PDF contracts, extracts the exact renewal dates into a database, and automatically triggers a 90-day warning email to the leasing team to start negotiations. Total automation."

Facilitator Note: 13:00 – 14:00 | LUNCH BREAK (60 min).

14:00 – 14:30 Workflow Automation Part 2 — Live Mini-Demo

Amardeep: "After lunch, one more live build — this time end to end, in front of you, using a problem this room actually raised over the last two days."

[Stage Direction] Build a second, shorter automation live (e.g. a document-intake-to-summary workflow) referencing a bottleneck a participant named during Day 2's AMA — this callback is deliberate.

14:30 – 15:30 The 1% Advantage — The Motivational Core

Amardeep: "We have reached the end of the technical curriculum. I want you all to take a deep breath. Look at how far we have come since Tuesday morning. You now understand Neural Networks, Transformers, Prompt Engineering, RLHF, APIs, and Custom GPTs. I can say with total confidence: you are officially in the top 1% of the global workforce when it comes to AI literacy. The other 99% are still typing 'write my email' into ChatGPT."

Amardeep: "But here is the harsh, absolute truth."

[Stage Direction] Amardeep walks to the center, stops pacing. Looks directly at the audience.

Amardeep: "What you have acquired over these three days is just a map. A map is beautiful. A map is accurate. But a map does not sail the ship. The foundation is done. Now, the real market separation begins. In the next 12 months, the global economy is going to brutally divide between the people who talk about AI in meetings, and the people who actually build with it. If you go back to your desk on Sunday morning and you do your job the exact same way you did it last week, these three days were a waste of your time."

😄 Amardeep (lighter, warmer, closing the loop): "Remember Ahmed from Accounts, who we've been laughing about since Tuesday morning? Here's the twist ending. The Ahmed we've been joking about all week isn't some hopeless case — he's every one of us before this workshop. And the version of Ahmed who walks out of this room today, builds his first automation, and shows up to work Monday morning already three steps ahead? That Ahmed just got the promotion. That's the whole story. That's why we're here." [Laugh Beat, then let it go quiet]

Facilitator Note: This is the full-circle payoff of the Ahmed running gag — three days of laughs land here as genuine motivation. Deliver it sincerely, not as a joke; the humor has already done its job of making the room fond of the character.

Facilitator Note: 15:30 – 15:45 | AFTERNOON BREAK (15 min).

15:45 – 16:15 The Pitch — The AI Practitioner Program

[Stage Direction] Click to Slide: Practitioner Program Cohorts.

Amardeep: "This exact gap — the gap between knowing the theory and building the systems — is why Webisdome and Traininglobe created the Practitioner Program. In this Foundation workshop, you sat and watched us explain the concepts. In the Practitioner Program, there are barely any slides. It is a builder's cohort. You open your laptops. You write the API connections. You build your own Custom GPTs using your company's actual, live data. You deploy real automated workflows before you even leave the room. Atharv will be looking over your shoulder, helping you debug your code."

Atharv: "If you want to be the person who brings AI into your company and gets the promotion because you saved them 500 hours of labor, the Practitioner Program is your blueprint."

Amardeep: "Spaces are strictly limited because we need small groups for that level of personalized technical attention. Because you committed to these three days, you are our Foundation Alumni. We reward action. If you go to the registration desk in the lobby today and secure your spot, you receive an on-the-spot 25% discount for Cohort 1 or Cohort 2."

[Stage Direction] Click to slide: Special Bonus / Merit Awards.

Amardeep: "And for our Top 5 Merit Award winners — the people who showed the best prompts, the best engagement over these three days — Webisdome believes in backing top talent. We are giving you a 30% Scholarship: we pay 30% of your tuition, you only pay the remaining 70%. We will read those names in just a moment."

Facilitator Note: Practitioner Program curriculum reference (for Q&A at the registration desk): Practical AI Use Cases · Advanced Prompting for Consistency · Structured Outputs Engineering · Knowledge-Based AI (RAG Basics) · Workflow Integration Design · Testing, Evaluation & Iteration · Live Project Day.

Facilitator Note: For anyone interested but not ready to commit today: they receive an invite to a free 60-90 minute Zoom livestream, "Practitioner AI Open Session" — a taste of RAG basics or structured outputs, exclusive to Foundation alumni. Mention this explicitly as a lower-commitment on-ramp.

16:15 – 16:45 Certificates & Final Send-Off

Amardeep: "Don't let the momentum die in this room. Don't go back to the old way. Go to the desk, claim your spot, and let's keep building the future. It has been our absolute honor to teach you. Thank you, Traininglobe. Thank you, SIG Oman. And thank you, Muscat!"

[Stage Direction] High-energy, upbeat, cinematic music starts playing instantly.

Amardeep: "Please join us at the photo booth. We are going to hand out the certificates and take photos with each of you!"

Facilitator Note: Photo booth backdrop per raw notes: frame featuring OpenAI, Claude, and their CEOs' likenesses with the line "I just completed my foundational AI course." Confirm image/likeness usage rights before printing this backdrop.

16:45 – 17:00 Certificate Ceremony, Photos & Registration Push

[Stage Direction] Amardeep and Atharv move immediately to the certificate table. They shake hands, take photos, and naturally, warmly direct enthusiastic participants straight toward the registration desk to secure the Practitioner Program upsell.

ALT. Day 3 — Option 1 (Alternate Running Order)

Facilitator Note: Use this sequence instead of the build above if you prefer a topic-by-topic structure with a dedicated video/poll/quiz per topic, matching the Day 1/Day 2 rhythm exactly. Each numbered topic below gets roughly 20-25 minutes including its video cue and interactive element, run in this order: (10) AI Strategy, (11) Functional AI, (12) Ethics and Governance in AI, (13) Cyber Security in the Age of AI, (14) Introduction to Other AI Tools, (15) Workflow Automation / Prompt Engineering. Follow with Awards & Recognition, Certificate Ceremony, LinkedIn Photoshoot, and the Practitioner Program pitch exactly as scripted above from 15:45 onward.

APPENDICES

Appendix A — Day 4 to Day 365: The Alumni Handholding Program

Facilitator Note: This is the long-tail retention and upsell structure referenced in the raw notes. It is not part of the live 3-day script but should be described briefly during the Day 3 close if time allows, or left to a follow-up email.

- Quarterly Refresher Course — a live session every quarter for all Foundation alumni, refreshing core concepts and introducing what has changed in the AI landscape since the last session.
- WhatsApp AI Foundation Alumni Group — a moderated, broadcast-style channel (admin-posted only, not an open chat) to keep noise low and signal high.
- Monthly "Ask Me Anything" Window — the first Monday of every month, one hour, open replies. Any course-related query can be raised during this window. Each monthly AMA is hosted by a different subject-matter expert, so alumni get exposure to a rotating set of voices, not just Amardeep and Atharv.
- The quarterly refresher remains the flagship event of the alumni program; the WhatsApp group and monthly AMA exist to sustain engagement between quarters.

Appendix B — AI Practitioner Program: Full Curriculum

1. Practical AI Use Cases
2. Advanced Prompting for Consistency
3. Structured Outputs Engineering
4. Knowledge-Based AI (RAG Basics)
5. Workflow Integration Design
6. Testing, Evaluation & Iteration
7. Live Project Day

Facilitator Note: On-ramp for the not-yet-committed: a free 60-90 minute Zoom "Practitioner AI Open Session," exclusive to Foundation alumni, previewing RAG basics or structured outputs.

Facilitator Note: Incentives: first 5 registrants receive a 30% scholarship (Webisdom covers 30%, participant pays 70%); the same or an overlapping top-5 group also receives a Certificate of Merit for engagement during the Foundation workshop. On-the-spot registration on Day 3 carries a 25% discount for Cohort 1 or Cohort 2.

Appendix C — Backup / Overflow Content Bank

Facilitator Note: Per the raw notes: prepare roughly 6 hours of extra material per day beyond the scripted 8, then trim in rehearsal. Use the items below as swap-ins if a segment runs short, or cuts if a segment runs long.

- A curated reel of "Best AI videos" — a rotating bank of short (60-90 second) demo clips (image generation, video generation, agentic tool use, robotics) to slot into any session that finishes early.
- Extra memes for each core concept (Intelligence Ladder, ML vs DL vs AI, ANI/AGI/ASI, Neural Networks, Transformers, RLHF) — one to two per topic, used only if the room's energy dips.
- Additional polls/quizzes/puzzles for every numbered topic across all three days — held in reserve so no session ever runs on empty air if a live demo fails or loads slowly.

Facilitator Note: Facilitator discipline: it is easier to cut prepared content live than to invent it live. Rehearse the full 6-hour surplus at least once so both facilitators know exactly what to drop first, second, and third if the clock runs tight.